

CDS503: Machine Learning

LAB 4: Naive Bayes (NB)


In this lab, we will be taking a *closer* look at **naive Bayes classification**.

Naive Bayes models are a group of extremely fast and simple classification algorithms that are often suitable for very high-dimensional datasets. Because they are so fast and have so few tunable parameters, they end up being very useful as a *quick-and-dirty* **baseline** for a classification problem. This lab will focus on an *intuitive* explanation of how naive Bayes classifiers **work**, followed by a couple examples of them in action on some datasets.

Step1: Business Understanding

This data was extracted from the census bureau database found at: <http://www.census.gov/ftp/pub/DES/www/welcome.html> (<http://www.census.gov/ftp/pub/DES/www/welcome.html>). It contains the census income of the people. They are trying to see the income of more than 50k and less than 50k. | Probability for the label '>50K' : 23.93% / 24.78% (without unknowns) | Probability for the label '<=50K' : 76.07% / 75.22% (without unknowns)

Step 1: Business Understanding




Census Bureau Data Set

Background:

- This data was extracted from the census bureau database found at: <https://archive.ics.uci.edu/ml/datasets/census+income>. It contains the census income of the people.
- They are trying to see the income of more than 50k and less than 50k. | Probability for the label '>50K' : 23.93% / 24.78% (without unknowns) | Probability for the label '<=50K' : 76.07% / 75.22% (without unknowns).
- There are two sets of data, one belongs for training and the other one is for testing.
- In training samples it contains 32560 samples and 15 attributes while for testing data it contains 16281 samples and 15 attributes.

Objectives

To do a classification on income of the people by using Naïve Bayes algorithm.



Bayesian Classification

Naive Bayes classifiers are built on Bayesian classification methods. These rely on **Bayes's theorem**, which is an equation describing the *relationship* of **conditional probabilities** of statistical quantities. In Bayesian classification, we're interested in finding the **probability** of a label *given* some **observed** features, which we can write as $P(L \mid \text{features})$. Bayes's theorem tells us how to express this in terms of quantities we can compute more directly:

$$P(L \mid \text{features}) = \frac{P(\text{features} \mid L)P(L)}{P(\text{features})}$$

If we are trying to decide between two labels—let's call them L_1 and L_2 —then one way to make this decision is to compute the ratio of the posterior probabilities for each label:

$$\frac{P(L_1 \mid \text{features})}{P(L_2 \mid \text{features})} = \frac{P(\text{features} \mid L_1) P(L_1)}{P(\text{features} \mid L_2) P(L_2)}$$


All we need now is some model by which we can compute $P(\text{features} \mid L_i)$ for each label. Such a model is called a *generative model* because it specifies the hypothetical random process that generates the data. Specifying this generative model for each label is the main piece of the training of such a Bayesian classifier. The general version of such a training step is a very difficult task, but we can make it simpler through the use of some simplifying assumptions about the form of this model.

What does it mean? For example, it means we have to assume that the comfort of the room on the Titanic is *independent* of the fare ticket. This **assumption** is absolutely *wrong* and it is why it is called **Naive**. It allows to simplify the calculation, even on very large datasets.

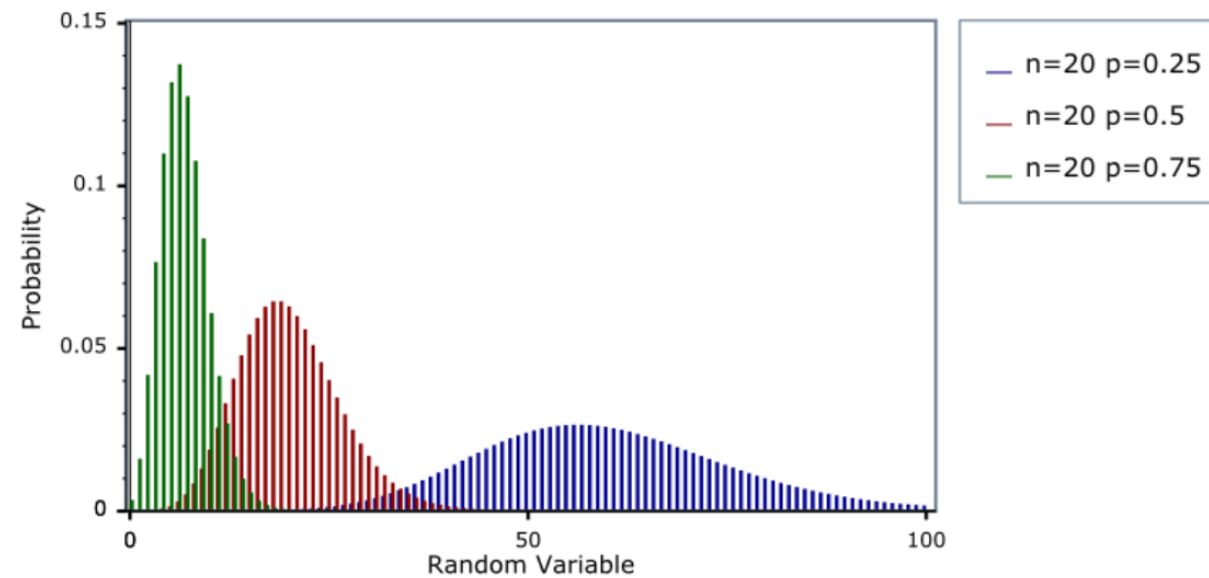
Different distribution function

To begin to implement a classifier, there is a few probability model *options* in the `sklearn` python library. They are:

- **Gaussian**: It *assumes* that **continuous features** follow a normal distribution.

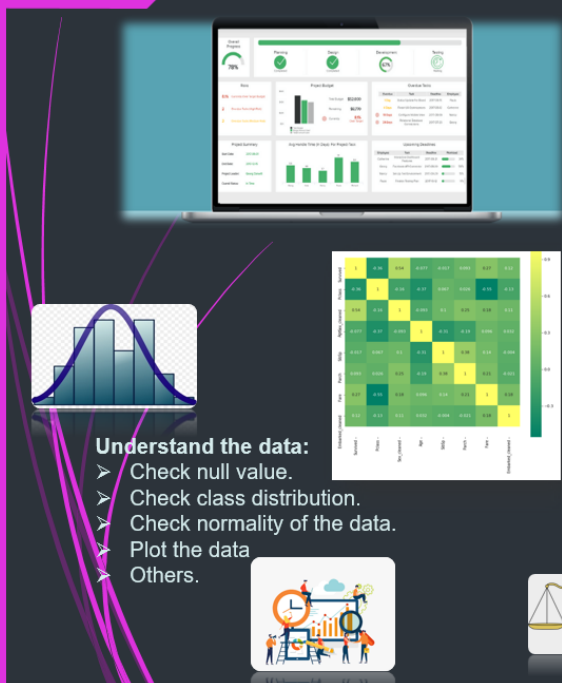


- **Multinomial**: It is useful if your **features are discrete**.
- **Bernoulli**: The binomial model is useful if your **features are binary**.



Different types of naive Bayes classifiers rest on **different naive assumptions** about the data, and we will examine a few of these in the following sections. Here we implement a classic **Gaussian Naive Bayes** on the **Titanic Disaster** dataset. We begin with the *standard library imports*.

Step 2: Data Understanding



Understand the data:

- Check null value.
- Check class distribution.
- Check normality of the data.
- Plot the data
- Others.

Description of the data:

In the census income data set, there are twelve attributes including the class attribute indicating the class/category information. The ten attributes are:

- age: continuous.
- workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
- fnlwgt: continuous.
- education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
- education-num: continuous.
- marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
- occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
- relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
- race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
- sex: Female, Male.
- capital-gain: continuous.
- capital-loss: continuous.
- hours-per-week: continuous.
- native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

Step 2: Data Understanding

Description of the data:

In the census income data set, there are fifteen attributes including the class attribute indicating the class/category information. The 15 attributes are:

- age: continuous.
- workclass:
 - Private
 - Self-emp-not-inc
 - Self-emp-inc
 - Federal-gov
 - Local-gov

- State-gov
- Without-pay
- Never-worked
- fnlwgt: continuous.
- education:
 - Bachelors
 - Some-college
 - 11th
 - HS-grad
 - Prof-school
 - Assoc-acdm
 - Assoc-voc
 - 9th
 - 7th-8th
 - 12th
 - Masters
 - 1st-4th
 - 10th
 - Doctorate
 - 5th-6th
 - Preschool
- education-num: continuous.
- marital-status:
 - Married-civ-spouse
 - Divorced
 - Never-married
 - Separated
 - Widowed
 - Married-spouse-absent
 - Married-AF-spouse
- occupation:
 - Tech-support
 - Craft-repair
 - Other-service
 - Sales

- Exec-managerial
- Prof-specialty
- Handlers-cleaners
- Machine-op-inspct
- Adm-clerical
- Farming-fishing
- Transport-moving
- Priv-house-serv
- Protective-serv
- Armed-Forces
- relationship:
 - Wife
 - Own-child
 - Husband
 - Not-in-family
 - Other-relative
 - Unmarried
- race:
 - White
 - Asian-Pac-Islander
 - Amer-Indian-Eskimo
 - Other
 - Black
- sex:
 - Female
 - Male
- capital-gain: continuous
- capital-loss: continuous
- hours-per-week: continuous.
- native-country:
 - United-States
 - Cambodia
 - England
 - Puerto-Rico
 - Canada

- Germany
- Outlying-US(Guam-USVI-etc)
- India
- Japan
- Greece
- South
- China
- Cuba
- Iran
- Honduras
- Philippines
- Italy
- Poland
- Jamaica
- Vietnam
- Mexico
- Portugal
- Ireland
- France
- Dominican-Republic
- Laos
- Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import time
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, precision_recall_fscore_support
%matplotlib inline
```



```
In [2]: # import library to display multiple outputs
from IPython.display import display

# Importing dataset
train = pd.read_csv("input/adult_train_modified.csv")
test = pd.read_csv("input/adult_test_modified.csv")

# see some of it, their overall statistics and dimensions
display(train.head(5))
display(train.describe())
display(train.shape)
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	Class
0	22	5	2491	9	12	4	0	1	4	1	24	0	39	38	0
1	33	4	2727	9	12	2	3	0	4	1	0	0	12	38	0
2	21	2	13188	11	8	0	5	1	4	1	0	0	39	38	0
3	36	2	14354	1	6	2	5	0	2	1	0	0	39	38	0
4	11	2	18120	9	12	2	9	5	2	0	0	0	39	4	0

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
count	30162.000000	30162.000000	30162.000000	30162.000000	30162.000000	30162.000000	30162.000000	30162.000000	30162.000000	30162.000000
mean	21.435482	2.199324	9825.221504	10.333764	9.121312	2.580134	5.959850	1.418341	3.678602	0.675685
std	13.125355	0.953925	5671.017927	3.812292	2.549995	1.498016	4.029566	1.601338	0.834709	0.468126
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	11.000000	2.000000	5025.250000	9.000000	8.000000	2.000000	2.000000	0.000000	4.000000	0.000000
50%	20.000000	2.000000	9689.500000	11.000000	9.000000	2.000000	6.000000	1.000000	4.000000	1.000000
75%	30.000000	2.000000	14520.750000	12.000000	12.000000	4.000000	9.000000	3.000000	4.000000	1.000000
max	71.000000	6.000000	20262.000000	15.000000	15.000000	6.000000	13.000000	5.000000	4.000000	1.000000

(30162, 15)

Step 3: Data Preparation

We will clean our data by replacing some missing values and encoding.

Step 3: Data Preparation

Definition: transformation of raw data into a form that is more suitable for modeling.

Why? :


- Machine learning algorithms require data to be numbers.
- Some machine learning algorithms impose requirements on the data.
- Statistical noise and errors in the data may need to be corrected.
- Complex nonlinear relationships may be teased out of the data.

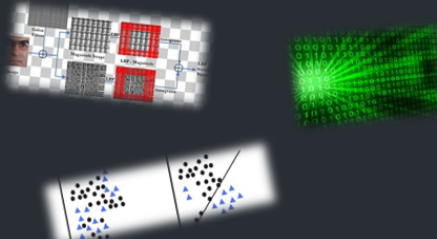
Tasks:

- Data Cleaning
- Feature Selection
- Data Transforms
- Feature Engineering
- Dimensionality Reduction

Other names:

- Data wrangling
- Data Cleaning
- Data pre-processing
- Feature Engineering





```

#Assign Column Name
colNames = ['age', 'workclass', 'fnlwt', 'education',
            'education-num', 'marital-status', 'occupation', 'relationship', 'race', 'sex', 'capital-gain',
            'capital-loss', 'hours-per-week', 'native-country', 'Class']

#Add different formats of missing value
missing_value_formats = ["n.a.", "?", "NA", "n/a", "na", "-.-"]

#read again training data
train = pd.read_csv('input/adult_train.csv', names=colNames, na_values = missing_value_formats, delimiter=' ', engine = 'python')
#read again test data
test = pd.read_csv('input/adult_test.csv', names=colNames, na_values = missing_value_formats, delimiter=' ', engine = 'python')

# see some of it, their overall statistics and dimensions
display(test.head(5))
display(test.describe())
display(test.shape)

```

In [3]: *#new sample after cleaning*
train.shape

Out[3]: (30162, 15)

```
In [4]: #new sample after cleaning  
test.shape
```

```
Out[4]: (15060, 15)
```

In [5]: *#Assign input column and label column for train data*

```
x_train = train.iloc[:, :-1]
```

```
y_train = train.iloc[:, 14]
```

```
print(x_train)
```

```
print(y_train)
```

	age	workclass	fnlwgt	education	education-num	marital-status	\
0	22	5	2491	9	12	4	
1	33	4	2727	9	12	2	
2	21	2	13188	11	8	0	
3	36	2	14354	1	6	2	
4	11	2	18120	9	12	2	
...	
30157	10	2	15471	7	11	2	
30158	23	2	7555	11	8	2	
30159	41	2	7377	11	8	6	
30160	5	2	12060	11	8	4	
30161	35	3	16689	11	8	2	

	occupation	relationship	race	sex	capital-gain	capital-loss	\
0	0	1	4	1	24	0	
1	3	0	4	1	0	0	
2	5	1	4	1	0	0	
3	5	0	2	1	0	0	
4	9	5	2	0	0	0	
...	
30157	12	5	4	0	0	0	
30158	6	0	4	1	0	0	
30159	0	4	4	0	0	0	
30160	0	3	4	1	0	0	
30161	3	5	4	0	107	0	

	hours-per-week	native-country
0	39	38
1	12	38
2	39	38
3	39	38
4	39	4
...
30157	37	38

30158	39	38
30159	39	38
30160	19	38
30161	39	38

[30162 rows x 14 columns]

0	0
1	0
2	0
3	0
4	0
	..
30157	0
30158	1
30159	0
30160	0
30161	1

Name: Class, Length: 30162, dtype: int64

In [6]: *#Assign input column and label column for test data*

```
x_test = test.iloc[:, :-1]
y_test = test.iloc[:, 14]
```

```
print(x_test)
print(y_test)
```

	age	workclass	fnlwgt	education	education-num	marital-status	\
0	8	2	8315	1	6	4	
1	21	2	1754	11	8	2	
2	11	1	10750	7	11	2	
3	27	2	4780	15	9	2	
4	17	2	7091	0	5	4	
...	
15055	16	2	8927	9	12	4	
15056	22	2	7893	9	12	0	
15057	21	2	11193	9	12	2	
15058	27	2	1593	9	12	0	
15059	18	3	6062	9	12	2	

	occupation	relationship	race	sex	capital-gain	capital-loss	\
0	6	3	2	1	0	0	
1	4	0	4	1	0	0	
2	10	0	4	1	0	0	
3	6	0	2	1	87	0	
4	7	1	4	1	0	0	
...	
15055	9	3	4	1	0	0	
15056	9	1	4	0	0	0	
15057	9	0	4	1	0	0	
15058	0	3	1	1	73	0	
15059	3	0	4	1	0	0	

	hours-per-week	native-country
0	39	37
1	49	37
2	39	37
3	39	37
4	29	37
...
15055	39	37

15056	35	37
15057	49	37
15058	39	37
15059	59	37

```
[15060 rows x 14 columns]
```

```
0      0
```

```
1      0
```

```
2      1
```

```
3      1
```

```
4      0
```

```
..
```

```
15055  0
```

```
15056  0
```

```
15057  0
```

```
15058  0
```

```
15059  1
```

```
Name: Class, Length: 15060, dtype: int64
```


Let's **clean** the data by *removing* the NaN data **completely**, and *replace* **categorical** values into **numeric** values:

Step 4: Modelling

Classify using Naive Bayes

Here, we separate the *descriptive* features and the *target* feature (`Income`) of the data. We assign the **training** set and **testing** set. After the label encoding, the income now becomes (`0` for `Income <=50k`) and (`1` for `income >50k`).

Step 4: Modelling



Machine Learning using Naïve Bayes

Naive Bayes classifiers are built on Bayesian classification methods. These rely on **Bayes's theorem**, which is an equation describing the *relationship* of **conditional probabilities** of statistical quantities. In Bayesian classification, we're interested in finding the **probability** of a label *given* some **observed** features, which we can write as $P(L \mid \text{features})$. Bayes's theorem tells us how to express this in terms of quantities we can compute more directly.

$$P(L \mid \text{features}) = \frac{P(\text{features} \mid L)P(L)}{P(\text{features})}$$

If we are trying to decide between two labels—let's call them L_1 and L_2 —then one way to make this decision is to compute the ratio of the posterior probabilities for each label:


$$\frac{P(L_1 \mid \text{features})}{P(L_2 \mid \text{features})} = \frac{P(\text{features} \mid L_1) P(L_1)}{P(\text{features} \mid L_2) P(L_2)}$$

All we need now is some model by which we can compute $P(\text{features} \mid L_i)$ for each label. Such a model is called a *generative model* because it specifies the hypothetical random process that generates the data. Specifying this generative model for each label is the main piece of the training of such a Bayesian classifier. The general version of such a training step is a very difficult task, but we can make it simpler through the use of some simplifying assumptions about the form of this model.

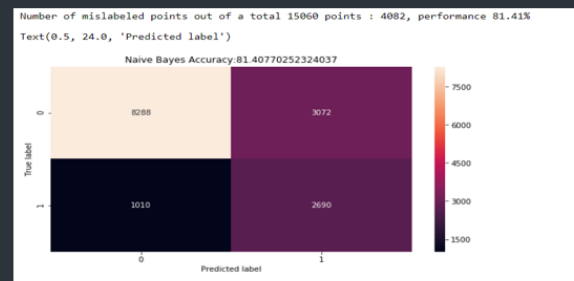

What does it mean? For example, it means we have to assume that the comfort of the room on the Titanic is *independent* of the fare ticket. This **assumption** is absolutely *wrong* and it is why it is called **Naïve**. It allows to simplify the calculation, even on very large datasets.

Step 5: Evaluation

Step 5: Evaluation



Confusion matrix



Accuracy Score

Then, we train the **Gaussian Naive Bayes** classifier and test it to see its *performance* in term of **accuracy**. From there, we can observe the misclassified ones using the *confusion matrix*.

- **Bernoulli Naive Bayes** : It assumes that all our features are binary such that they take only two values. Means 0s can represent "word does not occur in the document" and 1s as "word occurs in the document" .
- **Multinomial Naive Bayes** : Its is used when we have discrete data (e.g. movie ratings ranging 1 and 5 as each rating will have certain frequency to represent). In text learning we have the count of each word to predict the class or label.
- **Gaussian Naive Bayes** : Because of the assumption of the normal distribution, Gaussian Naive Bayes is used in cases when all our features are continuous. For example in Iris dataset features are sepal width, petal width, sepal length, petal length. So its features can have different values in data set as width and length can vary. We can't represent features in terms of their occurrences. This means data is continuous. Hence we use Gaussian Naive Bayes here.

```

In [7]: from sklearn.naive_bayes import GaussianNB, BernoulliNB, MultinomialNB
# Instantiate the classifier
gnb = GaussianNB()
#bnb = BernoulliNB()
#mnb = MultinomialNB()

# Train classifier
gnb.fit(x_train,y_train)
#bnb.fit(x_train,y_train)
#mnb.fit(x_train,y_train)

# Test the classifier
predict = gnb.predict(x_test)
#predict = bnb.predict(x_test)
#predict = mnb.predict(x_test)

# Print results
print("Number of mislabeled points out of a total {} points : {}, performance {:.2f}%"
      .format(x_test.shape[0], (y_test != predict).sum(),
              gnb.score(x_test,y_test)*100 ))

# Creates a confusion matrix
cm = confusion_matrix(y_test, predict)

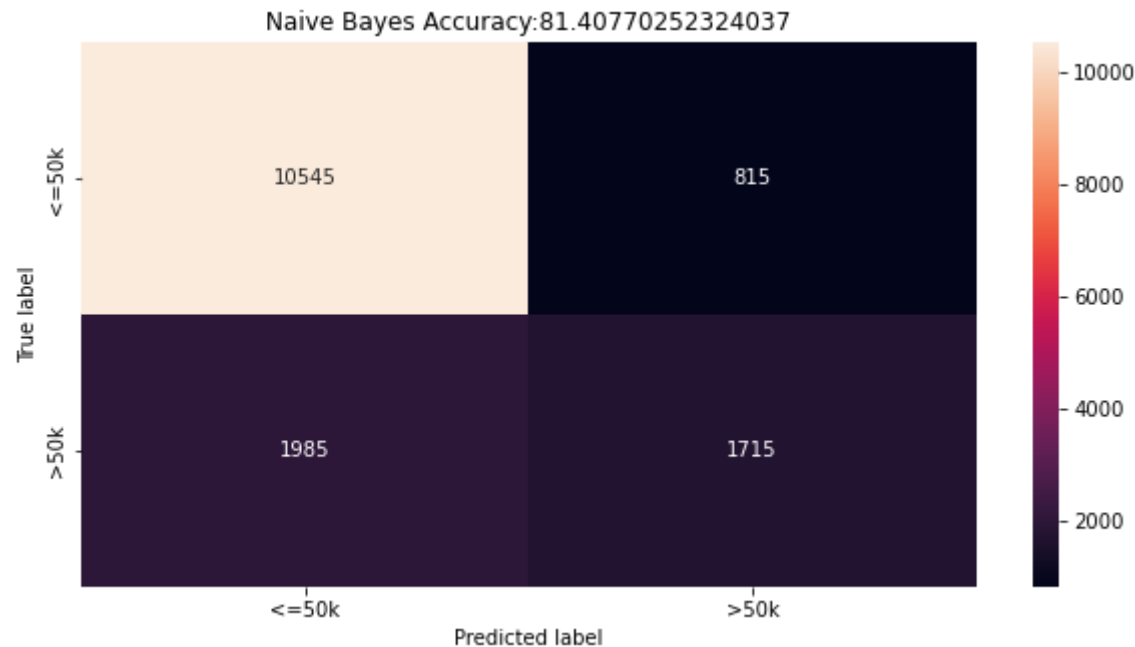
# Transform to dataframe for easier plotting
cm_df = pd.DataFrame(cm, index = ['<=50k', '>50k'],
                      columns = ['<=50k', '>50k'])

# plot the confusion matrix
plt.figure(figsize=(10,5))
sns.heatmap(cm_df, annot=True, fmt='g')
plt.title("Naive Bayes Accuracy:" + str(gnb.score(x_test,y_test)*100))
plt.ylabel('True label')
plt.xlabel('Predicted label')

```

Number of mislabeled points out of a total 15060 points : 2800, performance 81.41%

Out[7]: Text(0.5, 24.0, 'Predicted label')



Summary

Pros:

- Computationally **fast**
- Simple to implement
- Works well with **small datasets**
- Works well with **high dimensions**
- Perform well even if the Naive assumption is **not perfectly met**. In many cases, the approximation is enough to build a good classifier.

Cons:

- Require to **remove correlated features** because they are voted twice in the model and it can lead to over inflating importance (overfitting problem).
- If a **categorical variable** has a category in *testing* data set which was **not observed** in *training* data set, then the model will assign a **zero probability** and **failed** to make a *prediction*. This is often known as **"Zero Frequency"**. To solve this, we can use the **smoothing technique**.

One of the simplest smoothing techniques is called **Laplace estimation**. `sklearn` applies **Laplace smoothing** by default when you train a Naive Bayes classifier.

In []: